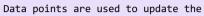


LLM-Elicited Priors

 $\theta \sim p_T(\Theta)$

The LLM outputs a Gaussian prior for each task description and feature, which we use to build a mixture model.



Data

prior to produce the posterior with Monte Carlo methods.

Calculate posterior

 $p(\theta|D) \propto p(D|\theta) p(\theta)$

Make predictions