

mixture model. $\theta \sim p_T(\Theta)$

LLM-Elicited Priors

The LLM outputs a Gaussian prior

for each task description and

feature, which we use to build a

Data points are used to update the prior to produce the posterior with Monte Carlo methods.

Calculate posterior

Make predictions

Data

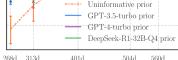
$p(\theta|D) \propto p(D|\theta) p(\theta)$







 $n = 5 \ n = 10$



n = 20

20 davs earlie